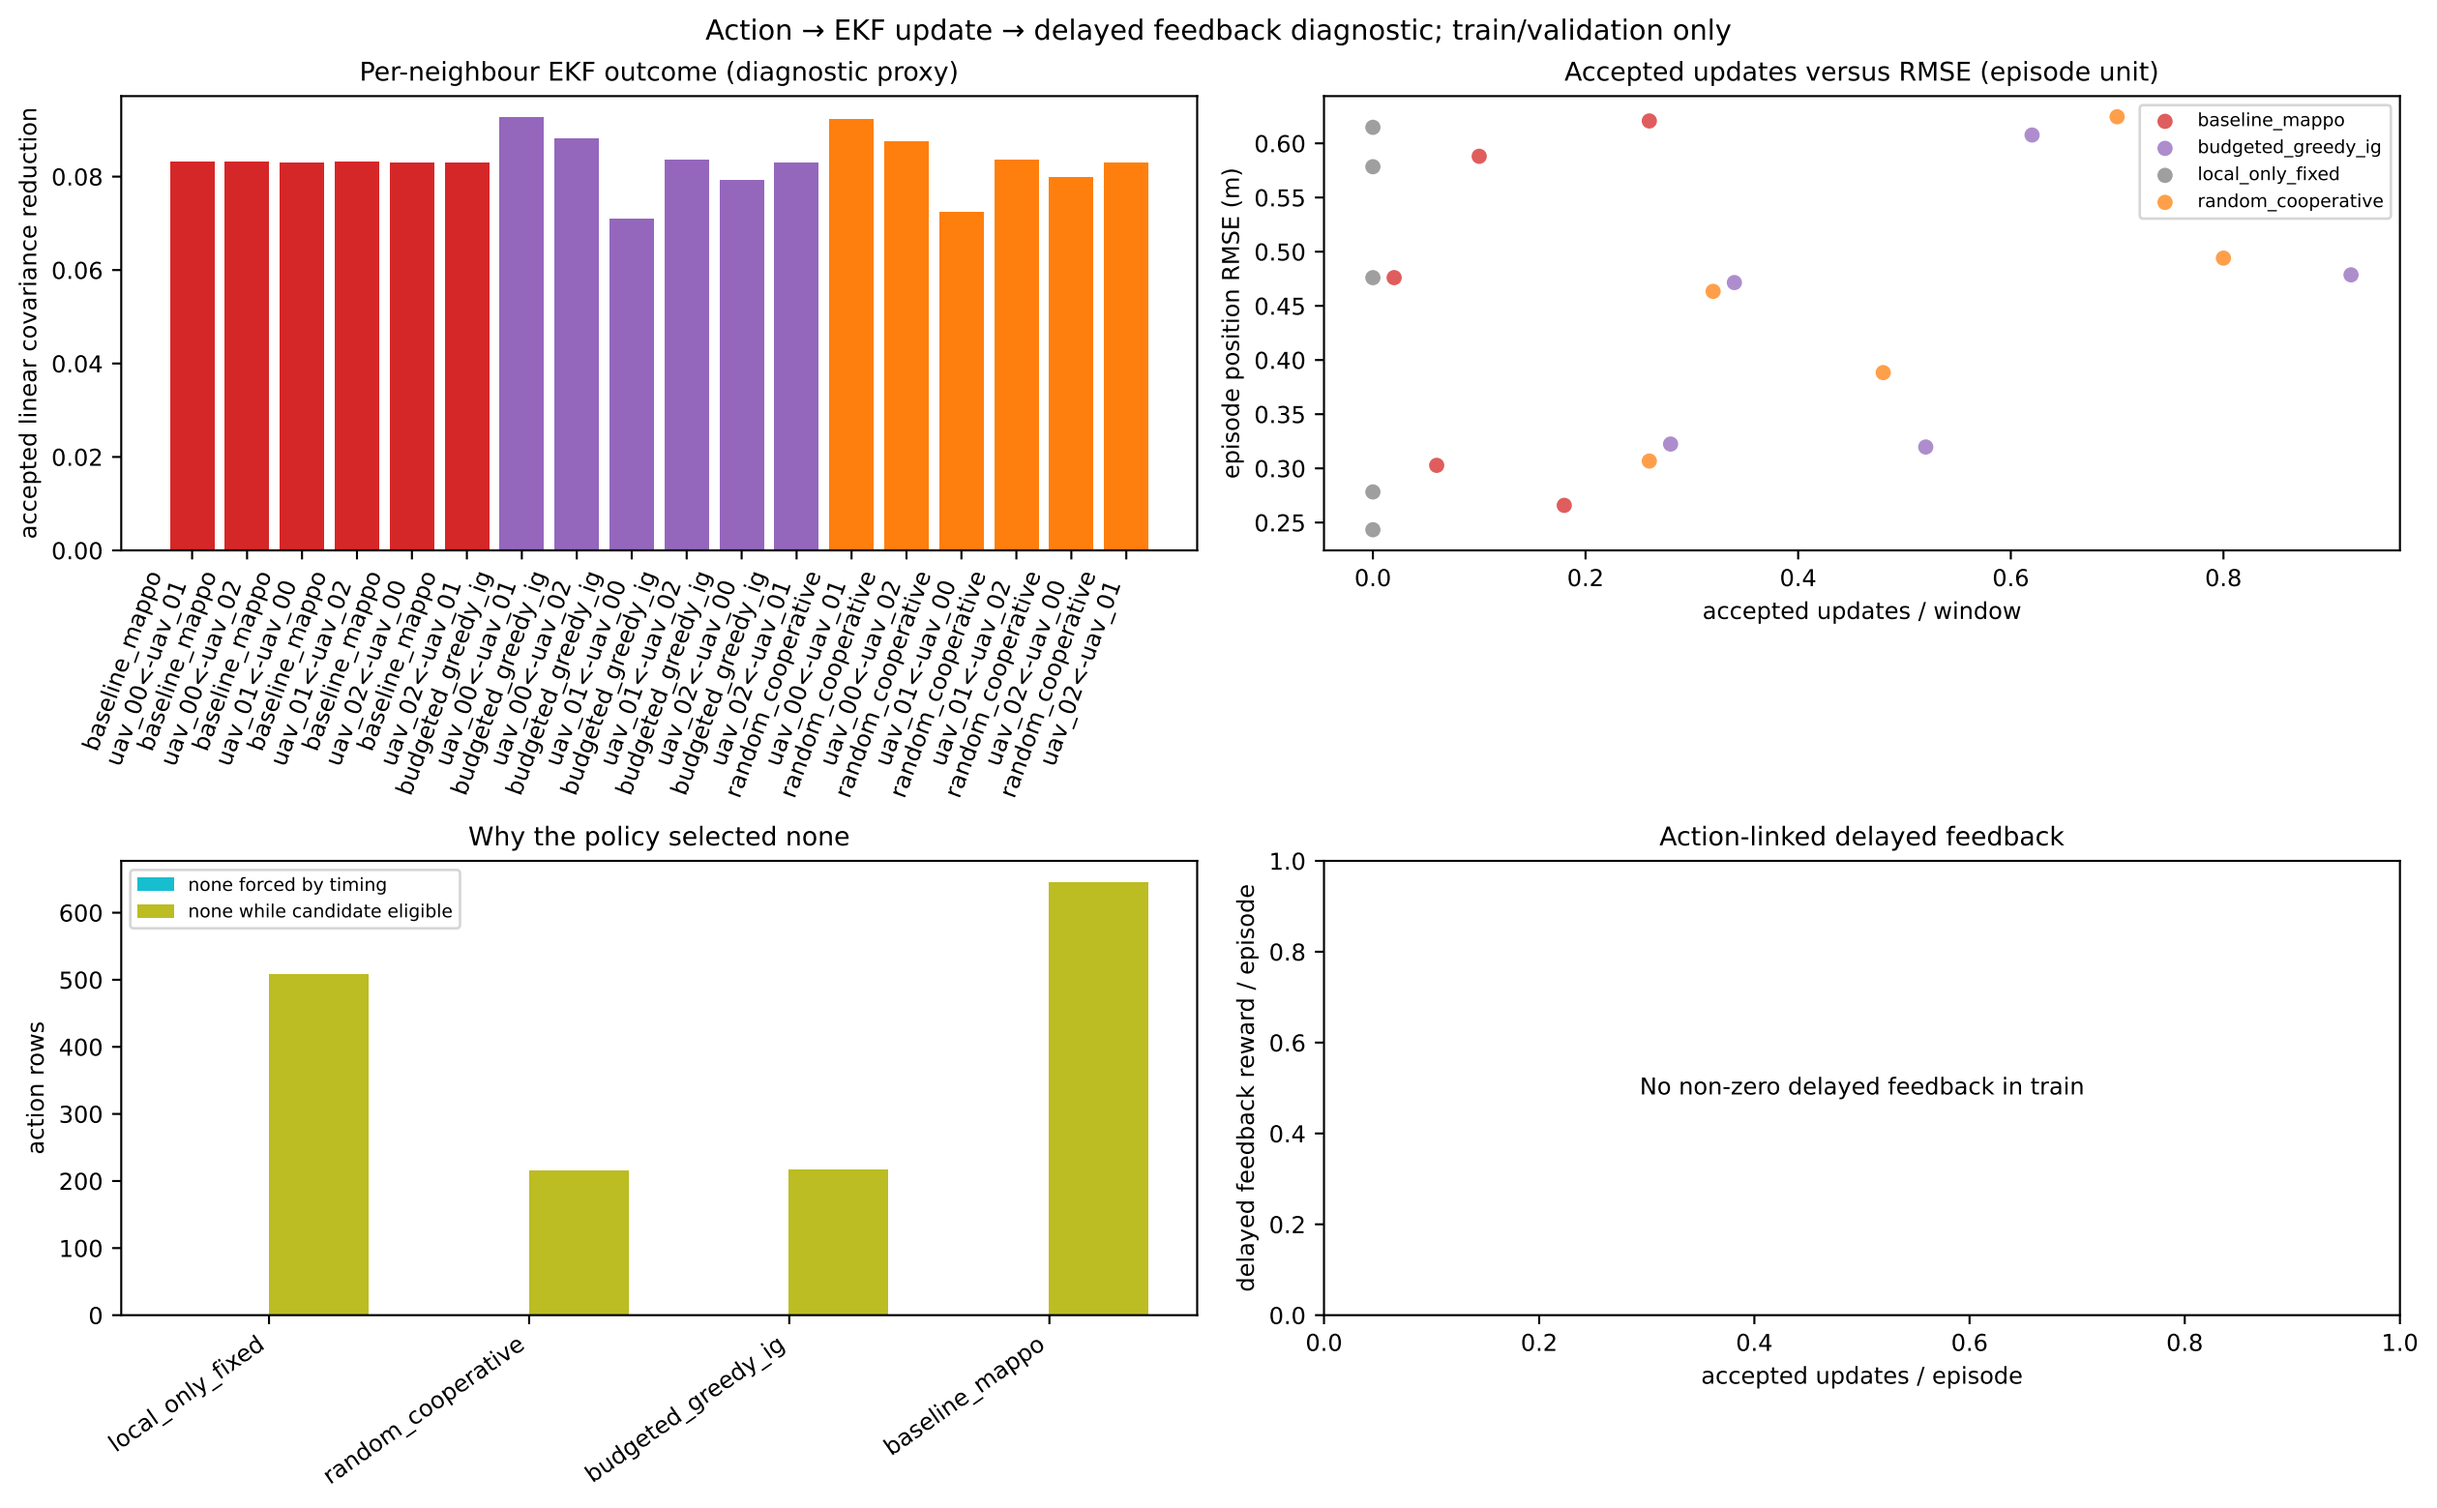
Action → EKF update → delayed feedback diagnostic; train/validation only
Per-neighbour EKF outcome (diagnostic proxy)
accepted linear covariance reduction
0.00
0.02
0.04
0.06
0.08
baseline_mappo uav_00<-uav_01
baseline_mappo uav_00<-uav_02
baseline_mappo uav_01<-uav_00
baseline_mappo uav_01<-uav_02
baseline_mappo uav_02<-uav_00
baseline_mappo uav_02<-uav_01
budgeted_greedy_ig uav_00<-uav_01
budgeted_greedy_ig uav_00<-uav_02
budgeted_greedy_ig uav_01<-uav_00
budgeted_greedy_ig uav_01<-uav_02
budgeted_greedy_ig uav_02<-uav_00
budgeted_greedy_ig uav_02<-uav_01
random_cooperative uav_00<-uav_01
random_cooperative uav_00<-uav_02
random_cooperative uav_01<-uav_00
random_cooperative uav_01<-uav_02
random_cooperative uav_02<-uav_00
random_cooperative uav_02<-uav_01
Accepted updates versus RMSE (episode unit)
episode position RMSE (m)
0.25
0.30
0.35
0.40
0.45
0.50
0.55
0.60
0.0
0.2
0.4
0.6
0.8
accepted updates / window
baseline_mappo
budgeted_greedy_ig
local_only_fixed
random_cooperative
Why the policy selected none
action rows
0
100
200
300
400
500
600
none forced by timing
none while candidate eligible
local_only_fixed
random_cooperative
budgeted_greedy_ig
baseline_mappo
Action-linked delayed feedback
delayed feedback reward / episode
0.0
0.2
0.4
0.6
0.8
1.0
No non-zero delayed feedback in train
accepted updates / episode